



Home



Training



Playground



Quick Training



Hands-On Labs



Learning Paths



Community



Search

← NGINX Web Server Deep Dive

Lecture: Load Balancing to Multiple Servers



Flash Cards

Notes

Transcript

Downloads

Looks like you haven't created any flash cards for this deck yet.

Create My First Flash Card

Popular Decks

Popular Cards

Nginx Web Server Deep Dive

NGINX Web Server Deep Dive

NGINX Web Server Deep Dive

In this lesson, we look at how we can utilize NGINX as a load balancer to distribute traffic between multiple instances of the same application.

Note: the commands in this video are run as the `root` user.

Documentation For This Video

- NGINX [http\\_upstream](#) module
- NGINX [upstream](#) directive
- NGINX [server](#) directive from the [upstream](#) module

Running Multiple Instances

To demonstrate load balancing, we need to run more than one instance of an application. We're going to create some duplicate services for `photos.example.com`, where the separate web-clients all run on different ports. To do this, we'll need to duplicate the `web-client.service` file a few times and add an additional `PORT` environment variable.

```
[root] $ cp /etc/systemd/system/web-client{,2}.service
[root] $ cp /etc/systemd/system/web-client{,3}.service
```

Now we set the `PORT` value in the new files:

`/etc/systemd/system/web-client2.service`

```
[Unit]
Description=S3 Photo App Node.js service
After=network.target photo-filter.target photo-storage.target

[Service]
Restart=always
User=nobody
Group=nobody
Environment=NODE_ENV=production
Environment=AWS_ACCESS_KEY_ID=YOUR_AWS_KEY_ID
Environment=AWS_SECRET_ACCESS_KEY=YOUR_AWS_SECRET_KEY
Environment=PORT=3100
ExecStart=/srv/www/s3photoapp/apps/web-client/bin/www

[Install]
WantedBy=multi-user.target
```

And for the third service:

`/etc/systemd/system/web-client3.service`

```
[Unit]
Description=S3 Photo App Node.js service
After=network.target photo-filter.target photo-storage.target

[Service]
Restart=always
User=nobody
Group=nobody
Environment=NODE_ENV=production
Environment=AWS_ACCESS_KEY_ID=YOUR_AWS_KEY_ID
Environment=AWS_SECRET_ACCESS_KEY=YOUR_AWS_SECRET_KEY
Environment=PORT=3101
ExecStart=/srv/www/s3photoapp/apps/web-client/bin/www

[Install]
WantedBy=multi-user.target
```

Lastly, we need to start these services:

```
[root] $ systemctl start web-client2
[root] $ systemctl start web-client3
```

Creating a Server Group

To load balance the traffic between our three identical services with NGINX we'll need to use the `http_upstream` module and the `upstream` directive. This directive creates a new context where we configure all of the load balancing behavior. Let's create our server group now:

`/etc/nginx/conf.d/photos.example.com`

```
upstream photos {
    server 127.0.0.1:3000;
    server 127.0.0.1:3100;
    server 127.0.0.1:3101;
}

server {
    listen 80;
    server_name photos.example.com;

    client_max_body_size 5m;

    location / {
        proxy_pass http://photos;
        proxy_http_version 1.1;
        proxy_set_header X-Forwarded-For $proxy_add_x_forwarded_for;
        proxy_set_header X-Real-IP $remote_addr;
        proxy_set_header Upgrade $http_upgrade;
        proxy_set_header Connection "upgrade";
    }

    location ~* \.(js|css|png|jpg|gif) {
        root /var/www/photos.example.com;
    }
}
```

Our `upstream` context is named `photos` and defines three servers. These servers use the `server` directive provided by the upstream module and their definition is similar to what we would provide to `proxy_pass` except that we will leave off the `http` or `https` scheme.

Besides setting up the `upstream` context, we also needed to change our `proxy_pass` to proxy traffic to the server group instead of a specific server, and we did this by replacing `127.0.0.1:3000` with the name of our server group `photos`.

If we reload NGINX and make 4 or 5 requests to `photos.example.com` we can check stdout from each service to see that they each received at least 1 request. Each will look something like this:

```
[root] $ systemctl status web-client3
? web-client3.service - S3 Photo App Node.js service
   Loaded: loaded (/etc/systemd/system/web-client3.service; disabled; vendor preset: disabled)
   Active: active (running) since Sun 2018-03-18 12:59:03 UTC; 15min ago
   Main PID: 1511 (node)
   CGroup: /system.slice/web-client3.service
           ??1511 node /srv/www/s3photoapp/apps/web-client/bin/www

Mar 18 12:59:03 keiththomps3.mylabserver.com systemd[1]: Started S3 Photo App Node.js service.
Mar 18 12:59:03 keiththomps3.mylabserver.com systemd[1]: Starting S3 Photo App Node.js service...
Mar 18 12:59:07 keiththomps3.mylabserver.com www[1511]: Listening on port 3101
Mar 18 13:10:34 keiththomps3.mylabserver.com www[1511]: GET /favicon.ico 404 118.663 ms - 150
Mar 18 13:14:05 keiththomps3.mylabserver.com www[1511]: GET / 200 413.474 ms - 1745
```



Exceeded my Expectations



Room for Improvement



Mark Complete & Start Next Lesson